

REPORT

- Navodita Mathur, Allie Azzarello, Naveena Nagaraju

Extract the files from the attached zip. Open terminal in the extracted folder and follow the steps in report and .txt file.

PART-1:

1. Install Cassandra on Ubuntu VMs:

- First, update the package index on each VM:
`sudo apt-get update`
- Install Java if it's not already installed:
`sudo apt-get install default-jdk`
- Add the Apache Cassandra repository keys:
`sudo curl -o /etc/apt/keyrings/apache-cassandra.asc
https://downloads.apache.org/cassandra/KEYS`
- Add the Apache Cassandra repository to the package sources list:
`echo "deb [signed-by=/etc/apt/keyrings/apache-cassandra.asc]
https://deb.debian.org/apache/cassandra/41x main" | sudo tee -a
/etc/apt/sources.list.d/cassandra.sources.list`
- Update the package index again:
`sudo apt-get update`
- Install Cassandra
`sudo apt-get install Cassandra`

2. Configure Cassandra nodes:

- Navigate to the Cassandra configuration directory:
`cd /etc/cassandra/`
 - Edit the `cassandra.yaml` file:
`sudo nano cassandra.yaml`
- Adjust the following settings:
- `listen_address`: Set the IP address for this node to listen on.
 - `rpc_address`: Set the IP address for remote procedure calls.
 - `seeds`: Add the IP addresses of all nodes in the cluster.
 - modify the following in the files:
`materialized_views_enabled: true`
`sasi_indexes_enabled: true`
`user_defined_functions_enabled: true`
`scripted_user_defined_functions_enabled: true`
 - modify the limits:
`read_request_timeout: 100000ms`
`range_request_timeout: 100000ms`

```
write_request_timeout: 20000ms
counter_write_request_timeout: 50000ms
cas_contention_timeout: 10000ms
truncate_request_timeout: 600000ms
request_timeout: 100000ms
slow_query_log_timeout: 50000ms
- Save and exit the file.
```

3. Start Cassandra nodes:

```
- Start Cassandra on each node:
    sudo service start cassandra
- Check the status to ensure it's running:
    sudo systemctl status cassandra
```

4. Verify cluster status:

```
- Check the status of the Cassandra cluster:
    nodetool status
- You should see all nodes listed and 'UN' (up and normal) status for each.
- Log into one of the nodes:
    cqlsh 10.254.0.26 9042 --request-timeout=60000
```

Output:

```
Mini-Project2 — ssh -i private_key ubuntu@155.98.37.64 — 135x24
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
-- Address      Load      Tokens   Owns (effective)  Host ID                               Rack
UN  10.254.0.26   141.12 KiB  16      63.3%             a8b59fcb-7ace-496e-902e-175ce8d880cf rack1
UN  10.254.3.185  141.13 KiB  16      64.3%             81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
UN  10.254.1.186  141.16 KiB  16      72.4%             78714d64-a079-49b9-bd80-b5815cc264f5 rack1

ubuntu@3:/ $ nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
-- Address      Load      Tokens   Owns (effective)  Host ID                               Rack
DN  10.254.0.26   -428476177 bytes  16      38.3%             a8b59fcb-7ace-496e-902e-175ce8d880cf rack1
UN  10.254.3.185  -2124761956 bytes  16      34.2%             81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
UN  10.254.1.186  -1304703570 bytes  16      27.5%             78714d64-a079-49b9-bd80-b5815cc264f5 rack1

ubuntu@3:/ $ nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
```

```
Mini-Project2 — ssh -i private_key ubuntu@155.98.37.64 — 135x24
UN 10.254.3.185 -5249387427 bytes 16 34.2% 81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
DN 10.254.1.186 -3644897644 bytes 16 27.5% 78714d64-a879-49b9-bd80-b5815cc264f5 rack1

ubuntu@3:~$ nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load
UN 10.254.0.26 -4654672612 bytes 16 38.3% a8b59fcb-7ace-496e-982e-175ce8d880cf rack1
UN 10.254.3.185 -5249387427 bytes 16 34.2% 81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
DN 10.254.1.186 -3644897644 bytes 16 27.5% 78714d64-a879-49b9-bd80-b5815cc264f5 rack1

ubuntu@3:~$ nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load Tokens Ows (effective) Host ID Rack
UN 10.254.0.26 -4654672612 bytes 16 38.3% a8b59fcb-7ace-496e-982e-175ce8d880cf rack1
UN 10.254.3.185 -5249387427 bytes 16 34.2% 81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
UN 10.254.1.186 -3767767283 bytes 16 27.5% 78714d64-a879-49b9-bd80-b5815cc264f5 rack1

ubuntu@3:~$ nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load Tokens Ows (effective) Host ID Rack
UN 10.254.0.26 -4654672612 bytes 16 38.3% a8b59fcb-7ace-496e-982e-175ce8d880cf rack1
UN 10.254.3.185 -5249387427 bytes 16 34.2% 81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
UN 10.254.1.186 -3767767283 bytes 16 27.5% 78714d64-a879-49b9-bd80-b5815cc264f5 rack1

ubuntu@3:~$ sudo systemctl status cassandra
cassandra.service - LSB: distributed storage system for structured data
Loaded: loaded (/etc/init.d/cassandra; generated)
Active: active (running) since Sun 2024-04-07 12:23:58 UTC; 7h ago
Docs: man:systemd-sysv-generator(8)
Process: 28645 ExecStop=/etc/init.d/cassandra stop (code=exited, status=0/SUCCESS)
Process: 28658 ExecStart=/etc/init.d/cassandra start (code=exited, status=0/SUCCESS)
Tasks: 216 (limit: 2301)
CGroup: /system.slice/cassandra.service
└─28755 /usr/bin/java -ea -ds:net.openhft... -XX:+UseThreadPriorities -XX:+HeapDumpOnOutOfMemoryError -
Apr 07 12:23:58 3-2 systemd[1]: Stopped LSB: distributed storage system for structured data.
Apr 07 12:23:58 3-2 systemd[1]: Starting LSB: distributed storage system for structured data...
Apr 07 12:23:58 3-2 systemd[1]: Started LSB: distributed storage system for structured data.
lines 1-13/13 (END)
```

```
Mini-Project2 — ssh -i private_key ubuntu@155.98.37.64 — 135x24
UN 10.254.3.185 -5249387427 bytes 16 34.2% 81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
DN 10.254.1.186 -3644897644 bytes 16 27.5% 78714d64-a879-49b9-bd80-b5815cc264f5 rack1

ubuntu@3:~$ nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load
UN 10.254.0.26 -4654672612 bytes 16 38.3% a8b59fcb-7ace-496e-982e-175ce8d880cf rack1
UN 10.254.3.185 -5249387427 bytes 16 34.2% 81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
DN 10.254.1.186 -3644897644 bytes 16 27.5% 78714d64-a879-49b9-bd80-b5815cc264f5 rack1

ubuntu@3:~$ nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load Tokens Ows (effective) Host ID Rack
UN 10.254.0.26 -4654672612 bytes 16 38.3% a8b59fcb-7ace-496e-982e-175ce8d880cf rack1
UN 10.254.3.185 -5249387427 bytes 16 34.2% 81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
UN 10.254.1.186 -3767767283 bytes 16 27.5% 78714d64-a879-49b9-bd80-b5815cc264f5 rack1

ubuntu@3:~$ nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load Tokens Ows (effective) Host ID Rack
UN 10.254.0.26 -4654672612 bytes 16 38.3% a8b59fcb-7ace-496e-982e-175ce8d880cf rack1
UN 10.254.3.185 -5249387427 bytes 16 34.2% 81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
UN 10.254.1.186 -3767767283 bytes 16 27.5% 78714d64-a879-49b9-bd80-b5815cc264f5 rack1

Last login: Sun Apr 7 17:23:25 2024 from 10.254.0.26
ubuntu@3:~$ nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load Tokens Ows (effective) Host ID Rack
UN 10.254.0.26 -4732821585 bytes 16 38.3% a8b59fcb-7ace-496e-982e-175ce8d880cf rack1
UN 10.254.3.185 -5319392744 bytes 16 34.2% 81153ddf-71ef-4a5a-bd45-b5e74995b467 rack1
UN 10.254.1.186 -3699981431 bytes 16 27.5% 78714d64-a879-49b9-bd80-b5815cc264f5 rack1

Output:
SELECT
WHERE

References:
1. Hadoop Singl
Rodrigo
Link - ht
e88c3d09a25f
2. https://hubd
3. https://medu

ubuntu@3:~$
```

PART-2:

1. Data preparation:

- Run the notebook to generate csv file "access_logs_comma.csv"
- Copy the csv file to VM

```
scp -i private_key access_logs_comma.csv ubuntu@<ip_address>:access_logs_comma.csv
```

2. Create a Keyspace and Table:

- Create a keyspace to contain our table:

```
CREATE KEYSPACE IF NOT EXISTS logs_keyspace WITH replication = {'class':  
'SimpleStrategy', 'replication_factor': 1};
```

- Create a table to store log data

```
CREATE TABLE IF NOT EXISTS logs_keyspace.logs_table (  
    log_id INT PRIMARY KEY,  
    ip_address TEXT,  
    path TEXT,  
    protocol TEXT,  
    request_method TEXT,  
    request_timestamp TIMESTAMP,  
    response_code TEXT,  
    size INT,  
    user_agent TEXT,  
);
```

3.. Import data into cassandra

Using CQL COPY command:

```
COPY logs_keyspace.logs_table(log_id, timestamp, ip_address, request_method, path,  
protocol, response_code, size, user_agent)  
FROM 'access_logs.csv'  
WITH HEADER = true ;
```

3. Verify Data Import:

- Connect to your Cassandra instance using cqlsh.

```
cqlsh 10.254.0.26 9042 --request-timeout=60000
```

- Query the data to ensure it has been imported successfully:

```
SELECT * FROM logs_keyspace.logs_table LIMIT 10;
```

```

... path TEXT,
...
    <identifier> <quotedName>
... protocol TEXT,
...
    <identifier> <quotedName>
... request_method TEXT,
...
    <identifier> <quotedName>
... request_timestamp TIMESTAMP,
...
    <identifier> <quotedName>
... response_code TEXT,
...
    <identifier> <quotedName>
... size INT,
...
    <identifier> <quotedName>
... user_agent TEXT,
... );
]
cqlsh> COPY logs_keyspace.logs_table(log_id, ip_address, path, protocol, request_method, request_timestamp, response_code, size, user_agent) FROM 'access_logs_comma.csv' WITH HEADER = true;
Using 1 child processes

Starting copy of logs_keyspace.logs_table with columns [log_id, ip_address, path, protocol, request_method, request_timestamp, response_code, size, user_agent].
Processed: 3236783 rows; Rate: 3596 rows/s; Avg. rate: 8334 rows/s
3236783 rows imported from 1 files in 0 day, 0 hour, 6 minutes, and 28.363 seconds (0 skipped).
cqlsh>

```

```

cqlsh> SELECT * FROM logs_keyspace.logs_table LIMIT 10;

log_id | ip_address | path | protocol | request_method | request_timestamp
-----+-----+-----+-----+-----+-----
response_code | size | user_agent

+-----+-----+-----+-----+-----+-----+
1792034 | 47.39.156.135 | /plugins/user/station/ | HTTP/1.1 | HEAD | 2022-04-01 00:00:00.000000+0000
404 | 0 | DirBuster-1.0-RC1 (http://www.owasp.org/index.php/Category:OWASP_DirBuster_Project)
302602 | 47.39.156.135 | /templates/system/images/Laptops.html | HTTP/1.1 | HEAD | 2022-04-02 00:00:00.000000+0000
404 | 0 | DirBuster-1.0-RC1 (http://www.owasp.org/index.php/Category:OWASP_DirBuster_Project)
2301876 | 47.39.156.135 | /plugins/system/legacy/content/tuning.css | HTTP/1.1 | HEAD | 2022-04-02 00:00:00.000000+0000
200 | 0 | DirBuster-1.0-RC1 (http://www.owasp.org/index.php/Category:OWASP_DirBuster_Project)
531141 | 96.32.128.5 | /images/M_images/198228_1.css | HTTP/1.1 | HEAD | 2022-04-04 00:00:00.000000+0000
404 | 0 | DirBuster-1.0-RC1 (http://www.owasp.org/index.php/Category:OWASP_DirBuster_Project)
2119753 | 47.39.156.135 | /images/stories/p2/ | HTTP/1.1 | HEAD | 2022-04-02 00:00:00.000000+0000
404 | 0 | DirBuster-1.0-RC1 (http://www.owasp.org/index.php/Category:OWASP_DirBuster_Project)
1416569 | 47.39.156.135 | /plugins/system/opa/ | HTTP/1.1 | HEAD | 2022-04-01 00:00:00.000000+0000
404 | 0 | DirBuster-1.0-RC1 (http://www.owasp.org/index.php/Category:OWASP_DirBuster_Project)
1817764 | 47.39.156.135 | /index2/20061211/ | HTTP/1.1 | GET | 2022-04-01 00:00:00.000000+0000
200 | 4302 | DirBuster-1.0-RC1 (http://www.owasp.org/index.php/Category:OWASP_DirBuster_Project)
693077 | 96.32.128.5 | /images/banners/holocaust.php | HTTP/1.1 | HEAD | 2022-04-04 00:00:00.000000+0000
404 | 0 | DirBuster-1.0-RC1 (http://www.owasp.org/index.php/Category:OWASP_DirBuster_Project)
2962449 | 82.209.218.4 | /apache-log/access.log | HTTP/1.1 | GET | 2022-02-24 00:00:00.000000+0000
206 | 1024 |
2333002 | 96.32.128.5 | /images/stories/slideshow/DivxToDVD.html | HTTP/1.1 | HEAD | 2022-04-04 00:00:00.000000+0000
404 | 0 | DirBuster-1.0-RC1 (http://www.owasp.org/index.php/Category:OWASP_DirBuster_Project)

(10 rows)
cqlsh>

```

PART-3

Create the following functions:

```

CREATE OR REPLACE FUNCTION logs_keyspace.state_group_count( state map<text, int>, type
text )
CALLED ON NULL INPUT
RETURNS map<text, int>
LANGUAGE java AS '

```

```
Integer count = (Integer) state.get(type); if (count == null) count = 1; else count++;  
state.put(type, count); return state; ' ;
```

```
CREATE OR REPLACE AGGREGATE logs_keyspace.group_count(text)  
SFUNC state_group_count  
SType map<text, int>  
INITCOND {};
```

```
CREATE CUSTOM INDEX user_agent_idx ON logs_keyspace.logs_table(user_agent) USING  
'org.apache.cassandra.index.sasi.SASIIndex' WITH OPTIONS = {'mode':'CONTAINS',  
'case_sensitive':'false'};
```

```
CREATE OR REPLACE FUNCTION logs_keyspace.state_group_and_count( state map<text, int>,  
type text )  
CALLED ON NULL INPUT  
RETURNS map<text, int>  
LANGUAGE java AS '  
Integer count = (Integer) state.get(type); if (count == null) count = 1; else count++;  
state.put(type, count); return state; ' ;
```

```
CREATE OR REPLACE FUNCTION logs_keyspace.state_max_group_count (state map<text, int>)  
CALLED ON NULL INPUT  
RETURNS text LANGUAGE JAVA AS '  
    String max_value = null;  
    int maxCount = 0;  
    String result = null;  
    for (String key : state.keySet()) {  
        int value = state.get(key);  
        if (value > maxCount) {  
            max_value = key; maxCount = value;  
        }  
    }  
    result = max_value + ": " + Integer.toString(maxCount);  
    return result; ' ;
```

```
CREATE OR REPLACE AGGREGATE logs_keyspace.max_group_count(text)  
SFUNC state_group_and_count  
SType map<text, int>  
FINALFUNC state_max_group_count  
INITCOND {};
```



```

CREATE OR REPLACE FUNCTION logs_keyspace.state_group_count_having( state map<text,
int>, type text )
CALLED ON NULL INPUT
RETURNS map<text, int>
LANGUAGE java AS '
Integer count = (Integer) state.get(type); if (count == null) count = 1; else count++;
state.put(type, count); return state; ' ;

CREATE OR REPLACE FUNCTION logs_keyspace.group_count_having(state map<text, int>)
CALLED ON NULL INPUT
RETURNS text LANGUAGE JAVA AS '
String result = "";
for (String key : state.keySet()) {
    int value = state.get(key);
    if (value > 10) {
        result += key + ": " + Integer.toString(value) + ";";
    }
}
return result; ' ;

CREATE OR REPLACE AGGREGATE logs_keyspace.max_group_count_having(text)
SFUNC state_group_count_having
STYPE map<text, int>
FINALFUNC group_count_having
INITCOND {};

```

1. How many hits were made to the website item “/administrator/index.php”?

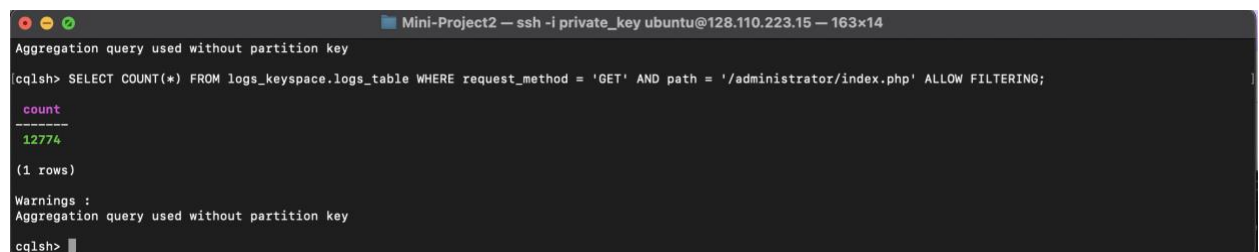
Query :

```

SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_method = 'GET' AND
path = '/administrator/index.php' ALLOW FILTERING;

```

Output:



```

Mini-Project2 — ssh -i private_key ubuntu@128.110.223.15 — 163x14
Aggregation query used without partition key
cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_method = 'GET' AND path = '/administrator/index.php' ALLOW FILTERING;
count
-----
12774
(1 rows)
Warnings :
Aggregation query used without partition key
cqlsh>

```

```
Mini-Project2 — ssh -i private_key ubuntu@155.98.37.64 — 130x22
ERING;
NoHostAvailable: ('Unable to complete the operation against any hosts', {<Host: 10.254.0.26:9042 datacenter1>: Unavailable('Error
from server: code=1000 [Unavailable exception] message="Cannot achieve consistency level ONE" info={\'consistency\': \'ONE\', \'re
quired_replicas\': 1, \'alive_replicas\': 0}\',)})})
cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_method = 'GET' AND path = '/administrator/index.php' ALLOW FILTERING;
ERING;
NoHostAvailable: ('Unable to complete the operation against any hosts', {<Host: 10.254.0.26:9042 datacenter1>: Unavailable('Error
from server: code=1000 [Unavailable exception] message="Cannot achieve consistency level ONE" info={\'consistency\': \'ONE\', \'re
quired_replicas\': 1, \'alive_replicas\': 0}\',)})})
cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_method = 'GET' AND path = '/administrator/index.php' ALLOW FILTERING;
ERING;

count
-----
12774

(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh>
```

2. How many hits were made from the IP: 96.32.128.5

Query:

```
SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE ip_address = '96.32.128.5'
ALLOW FILTERING;
```

Output:

```
Mini-Project2 — ssh -i private_key ubuntu@128.110.223.15 — 163x14
Aggregation query used without partition key

cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE ip_address = '96.32.128.5' ALLOW FILTERING;

count
-----
653428

(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh>
```

```
Mini-Project2 — ssh -i private_key ubuntu@155.98.37.64 — 130x22

count
-----
12774

(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE ip_address = '96.32.128.5' ALLOW FILTERING;

count
-----
653439

(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh>
```


3. Which path in the website has been hit most? How many hits were made to the path?

Query:

```
SELECT logs_keyspace.group_count(path) FROM logs_keyspace.logs_table;
```

Output:

4. Which IP accesses the website most? How many accesses were made by it?

Query:

```
SELECT logs_keyspace.group_count(ip_address) FROM logs_keyspace.logs_table;
```

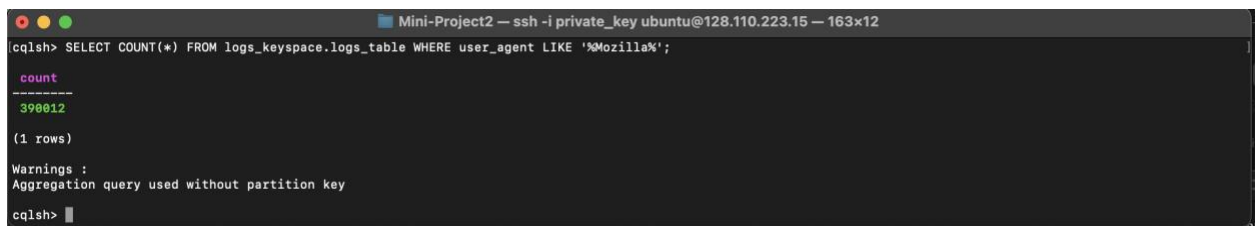
Output:

5. How many accesses were made by Firefox(Mozilla)?

Query:

```
SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE user_agent LIKE '%Mozilla%';
```

Output:

A terminal window titled "Mini-Project2 — ssh -i private_key ubuntu@128.110.223.15 — 163x12" displays the output of a MySQL query. The prompt is "cqlsh>". The query executed is "SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE user_agent LIKE '%Mozilla%';". The output shows a table with one column named "count" and one row with the value "390012". Below the table, it says "(1 rows)". A warning message follows: "Warnings : Aggregation query used without partition key". The prompt "cqlsh>" is shown again at the bottom.

```
Mini-Project2 — ssh -i private_key ubuntu@128.110.223.15 — 163x12
cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE user_agent LIKE '%Mozilla%';

count
-----
390012
(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh>
```

```
Mini-Project2 — ssh -i private_key ubuntu@155.98.37.64 — 130x22

count
-----
653439
(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE user_agent LIKE '%Mozilla%';

count
-----
390572
(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh>
```

6. For all requests on 02/Apr/2022, what is the ratio of GET request?

Query :

```
SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_timestamp >= '2022-04-02T00:00:00' AND request_timestamp < '2022-04-03T00:00:00' AND request_method = 'GET' ALLOW FILTERING;
```

```
Mini-Project2 — ssh -i private_key ubuntu@128.110.223.15 — 163x14

s." info={'consistency': 'ONE', 'required_responses': 1, 'received_responses': 0}
cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_timestamp >= '2022-04-02T00:00:00' AND request_timestamp < '2022-04-03T00:00:00' AND request_method = 'GET' ALLOW FILTERING;

count
-----
69102
(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh>
```

```
Mini-Project2 — ssh -i private_key ubuntu@155.98.37.64 — 130x22

count
-----
390572
(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_timestamp >= '2022-04-02T00:00:00' AND request_timestamp < '2022-04-03T00:00:00' AND request_method = 'GET' ALLOW FILTERING;

count
-----
69103
(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh>
```

```
SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_timestamp >= '2022-04-02T00:00:00' AND request_timestamp < '2022-04-03T00:00:00' ALLOW FILTERING;
```

Output:

```
Mini-Project2 -- ssh -i private_key ubuntu@128.110.223.15 -- 163x14
cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_timestamp >= '2022-04-02T00:00:00' AND request_timestamp < '2022-04-03T00:00:00' ALLOW FILTERING;

count
-----
442978
(1 rows)

Warnings :
Aggregation query used without partition key
cqlsh>
```

```
Mini-Project2 -- ssh -i private_key ubuntu@155.98.37.64 -- 130x22

count
-----
69103
(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_timestamp >= '2022-04-02T00:00:00' AND request_timestamp < '2022-04-03T00:00:00' ALLOW FILTERING;

count
-----
442985
(1 rows)

Warnings :
Aggregation query used without partition key
cqlsh>
```

So, $69102/442978 = 1/6$

```
Mini-Project2 -- ssh -i private_key ubuntu@128.110.223.15 -- 163x32
cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_timestamp >= '2022-04-02T00:00:00' AND request_timestamp < '2022-04-03T00:00:00' AND request_method = 'GET' ALLOW FILTERING;

count
-----
69102
(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE request_timestamp >= '2022-04-02T00:00:00' AND request_timestamp < '2022-04-03T00:00:00' ALLOW FILTERING;

count
-----
442978
(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh> SELECT 442978/69102 FROM system.local;

442978 / 69102
-----
6
(1 rows)
cqlsh>
```

```
Mini-Project2 — ssh -i private_key ubuntu@155.98.37.64 — 130x22

Warnings :
Aggregation query used without partition key

cqlsh> SELECT 442985.0/69103.0 FROM system.local
... ;
SyntaxException: line 1:0 no viable alternative at input 'SELECT' ([SELECT]...)
cqlsh> SELECT 442985.0/69103.0 FROM system.local ;

442985.0 / 69103.0
-----
6.4105

(1 rows)
cqlsh> SELECT 69103.0/442985.0 FROM system.local ;

69103.0 / 442985.0
-----
0.155994

(1 rows)
cqlsh>
```

7. How many requests are lower than or equal to 404 bytes?

Query:

```
SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE size <= 404 ALLOW FILTERING;
```

Output

```
Mini-Project2 — ssh -i private_key ubuntu@128.110.223.15 — 163x12

cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE size <= 404 ALLOW FILTERING;

count
-----
2716808

(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh>
```

```
Mini-Project2 — ssh -i private_key ubuntu@155.98.37.64 — 130x22

6.4105

(1 rows)
cqlsh> SELECT 69103.0/442985.0 FROM system.local ;

69103.0 / 442985.0
-----
0.155994

(1 rows)
cqlsh> SELECT COUNT(*) FROM logs_keyspace.logs_table WHERE size <= 404 ALLOW FILTERING;

count
-----
2716872

(1 rows)

Warnings :
Aggregation query used without partition key

cqlsh>
```

8. List the IPs that have more than ten 404 requests. If no ip fulfills, print the ip that has most 404 requests and the number of requests

Query:

```
SELECT logs_keyspace.max_group_count_having(ip_address) FROM  
logs_keyspace.logs_table WHERE response_code='404'ALLOW FILTERING;
```

Output:

If no IP fulfills the condition, then:

```
SELECT logs_keyspace.max_group_count(ip_address) FROM logs_keyspace.logs_table  
WHERE response_code='404'ALLOW FILTERING;
```

Output:

References:

1. Hadoop Single Node Cluster on Docker

Rodrigo Ancavil

Link - <https://medium.com/analytics-vidhya/hadoop-single-node-cluster-on-docker-e88c3d09a256>

2. https://hub.docker.com/_/eclipse-temurin

3. <https://medium.com/@abhikdey06/apache-hadoop-3-3-6-installation-on-ubuntu-22-04-14516bceec85>

4. <https://github.com/amephraim/nlp> (input text)